



# InVERGe: Intelligent Visual Encoder for Bridging Modalities in Report Generation

Reading Group  
09.06.205



**Input**



“stable mild cardiomegaly,  
no pneumothorax, pleural  
effusion, or focal airspace  
disease. bony structures  
intact. right humeral head  
bone anchor.”

**Output**

The main objective of this paper is to create clear reports about the image's contents.



Input

Encoder

Decoder

“stable mild cardiomegaly, no pneumothorax, pleural effusion, or focal airspace disease. bony structures intact. right humeral head bone anchor.”

Output

Trained on plenty of image-caption pairs

The main objective of this paper is to create clear reports about the image's contents.

Problem	Objective
<p>Models have been observed to produce <b>unnatural language outputs</b> when trained solely on raw image-text pairs</p>	<p>Introduce a <b>two-step training methodology</b> using a layer that bridges the gap between the encoder and decoder to significantly:</p> <ul style="list-style-type: none"> <li>a) <b>enhance the alignment</b> between vision and language</li> <li>b) <b>reduce the workload</b> on the decoder</li> </ul>
<p><b>High computational demand</b> of end-to-end model training</p>	

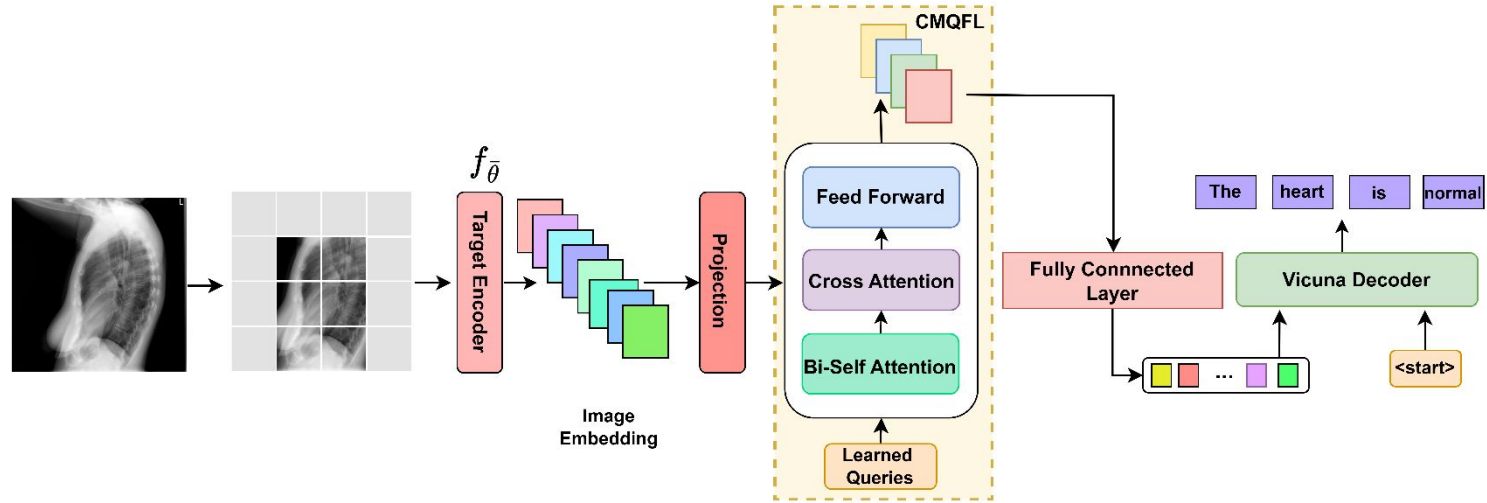
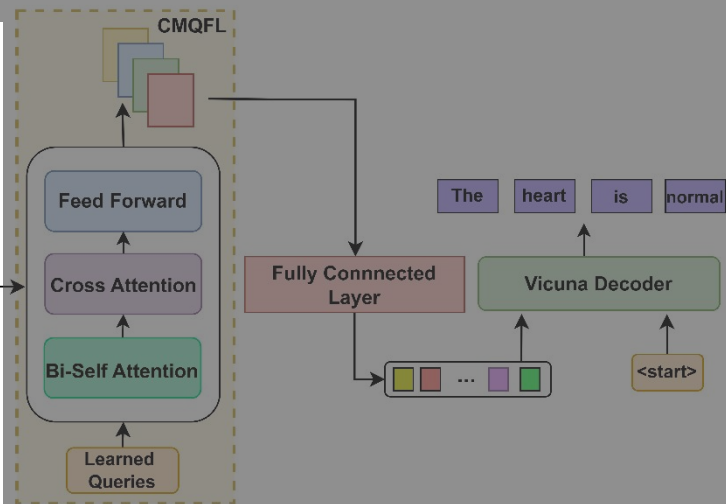

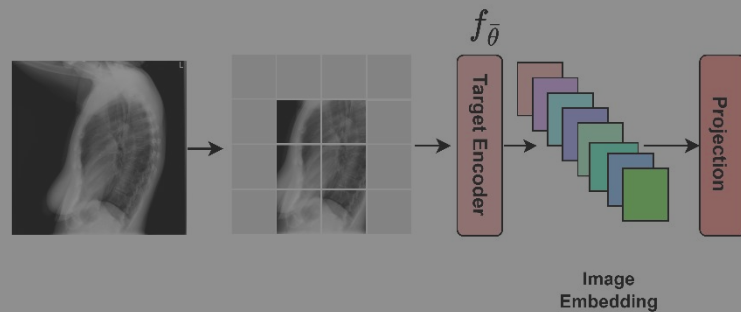


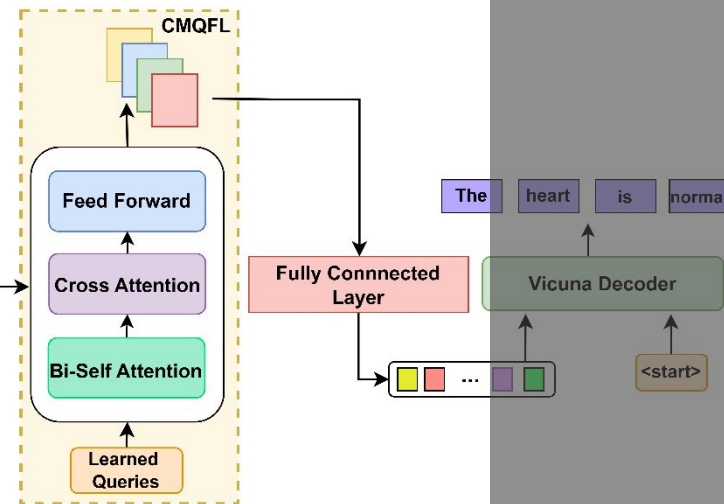
Figure 3. The architecture of the proposed InVERGe approach, includes our encoder trained on self-supervised methods, learnable CMQFL layers, and Vicuna decoder. Specifically, the encoder encodes high-level semantic features and the CMQFL layer uses that information and applies a query that will generate text-grounded image embeddings. Then the decoder uses that embedding to produce a report.

## (1) I-JEPA Image Encoder

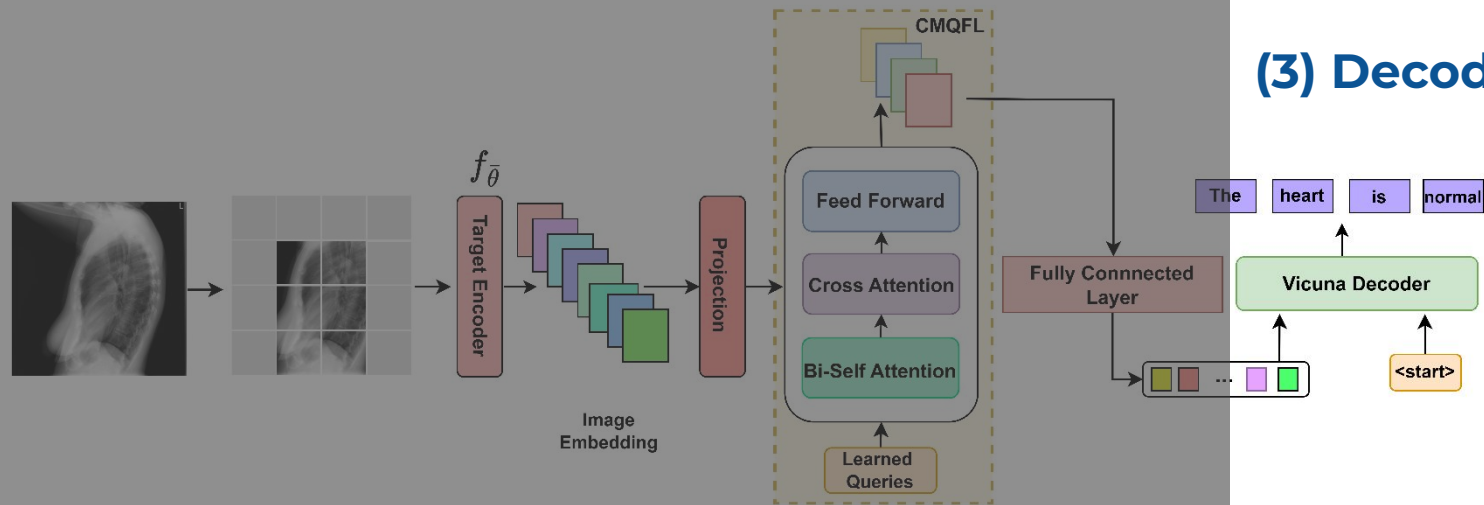




## (2) Cross-Modal Query Fusion Layer



Bert-based CMQFL layer  
responsible for generating  
text-grounded image  
embeddings



### (3) Decoder

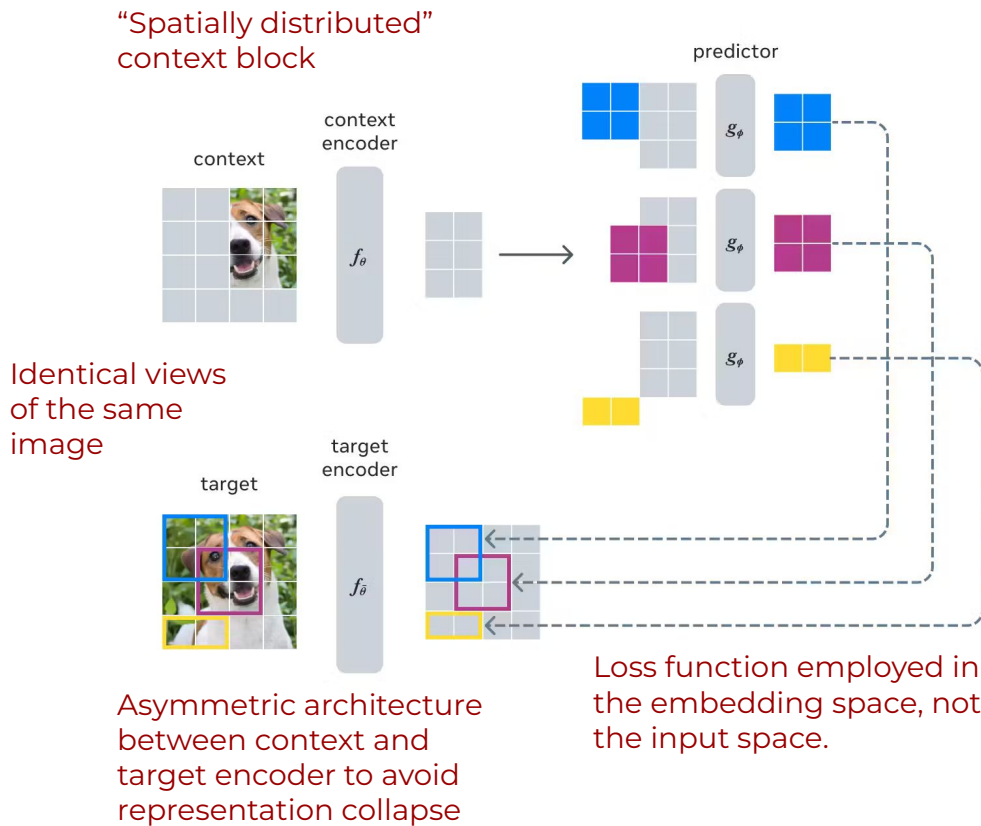


# Image Encoder

1/3

# What is JEPA?

- Joint Embedding Predictive Architecture
- Predict missing information in an **abstract representation space**



Why select **JEPA** as pre-training strategy?


## Why select **JEPA** as pre-training strategy?

Because of perceived issues with other pre-training strategies. Let's take a look at some:

# Why select **JEPA** as pre-training strategy?

Because of perceived issues with other pre-training strategies. Let's take a look at some:


## Invariance-based Pretraining



# Why select **JEPA** as pre-training strategy?

Because of perceived issues with other pre-training strategies. Let's take a look at some:

## Invariance-based Pretraining




Invariance-based pretraining methods optimize an encoder to produce similar embeddings for two or more views of the same image [15, 20], with image views typically constructed using a set of hand-crafted data augmentations, such as random scaling, cropping, and color jittering [20], amongst others [35]. These pretraining methods can produce representations of a high semantic level [4, 18], but they also introduce strong biases that may be detrimental for certain downstream tasks or even for pretraining tasks with different data distributions [2]. Often, it is unclear how to generalize these biases for tasks requiring different levels of abstraction. For example, image classification and instance segmentation do not require the same invariances [11]. Additionally, it is not straightforward to generalize these image-specific augmentations to other modalities such as audio.

# Why select **JEPA** as pre-training strategy?

Because of perceived issues with other pre-training strategies. Let's take a look at some:

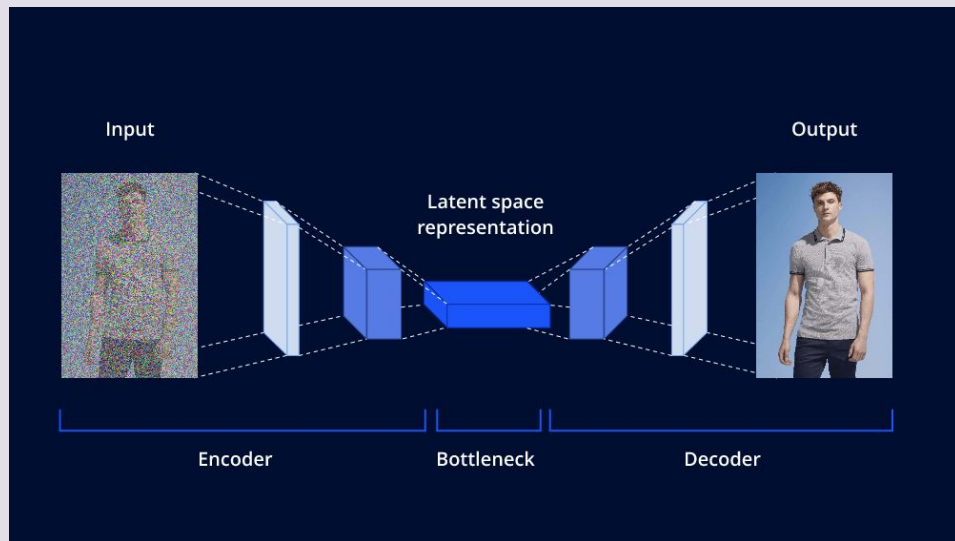
## Generative methods that predict in pixel/token space



# Why select **JEPA** as pre-training strategy?

Because of perceived issues with other pre-training strategies. Let's take a look at some:

## Generative methods that predict in pixel/token space



Loss function here is reconstruction error (in input space)

Compared to generative methods that predict in pixel/token space, I-JEPA makes use of abstract prediction targets for which unnecessary pixel-level details are potentially eliminated, thereby leading the model to learn more semantic features. Another core design choice to guide



How does InVERGe use JEPa?

## How does InVERGe use JEPA?

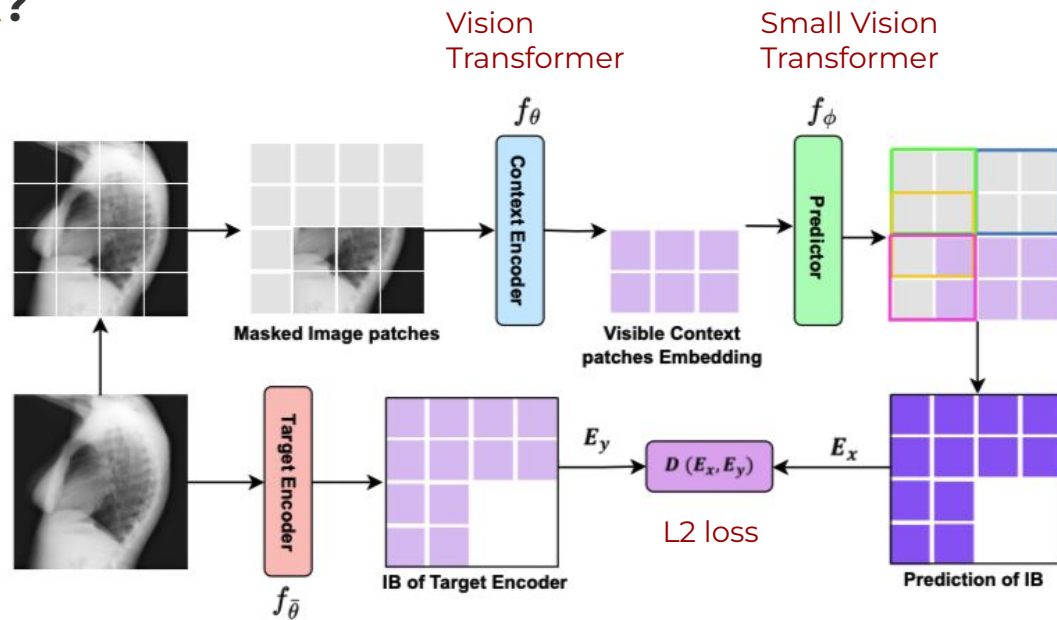


Figure 2. Visual representation of the initial step of training the encoder component of the InVERGe model.

## How does InVERGe use JEPA?

In this scenario, we employ gradient methods to fine-tune the parameters of both the predictor  $f_\phi$  and the CE ( $f_\theta$ ). Simultaneously, the parameters of TE ( $f_{\bar{\theta}}$ ) are continuously adjusted, achieved by applying an exponential moving average (EMA) technique to the parameters of the context encoder. Adopting an EMA strategy for target encoders is essential in achieving effective training results for joint embedding architectures (JEA) that incorporate ViT, as discussed in [5]. These updates also ensure the target

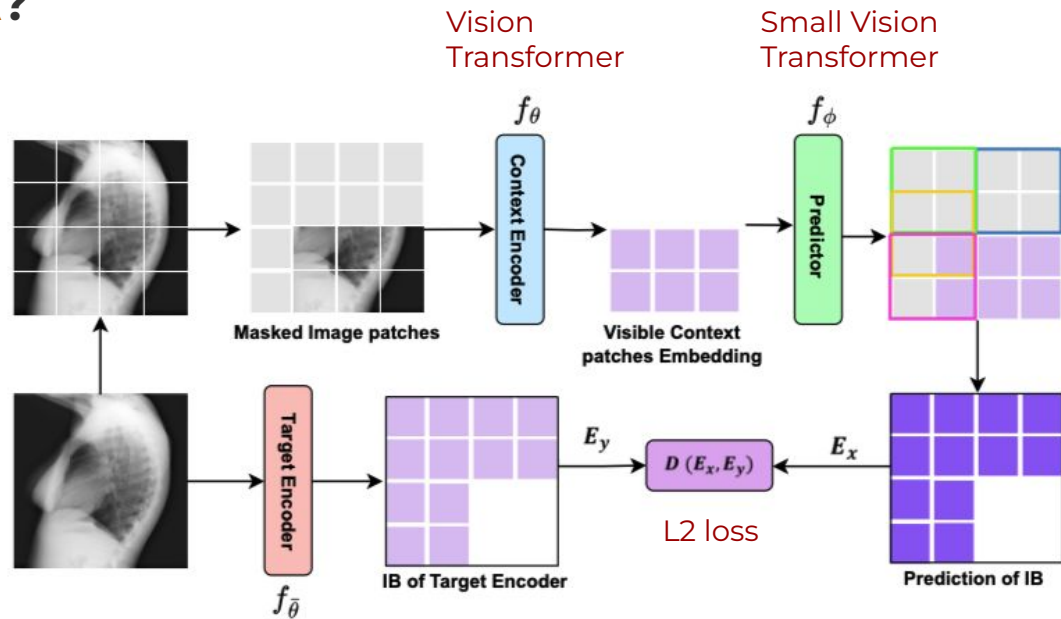


Figure 2. Visual representation of the initial step of training the encoder component of the InVERGe model.

## How does InVERGe use JEPA?

In this scenario, we employ gradient methods to fine-tune the parameters of both the predictor  $f_\phi$  and the CE ( $f_\theta$ ). Simultaneously, the parameters of TE ( $f_{\bar{\theta}}$ ) are continuously adjusted, achieved by applying an exponential moving average (EMA) technique to the parameters of the context encoder. Adopting an EMA strategy for target encoders is essential in achieving effective training results for joint embedding architectures (JEA) that incorporate ViT, as discussed in [5]. These updates also ensure the target

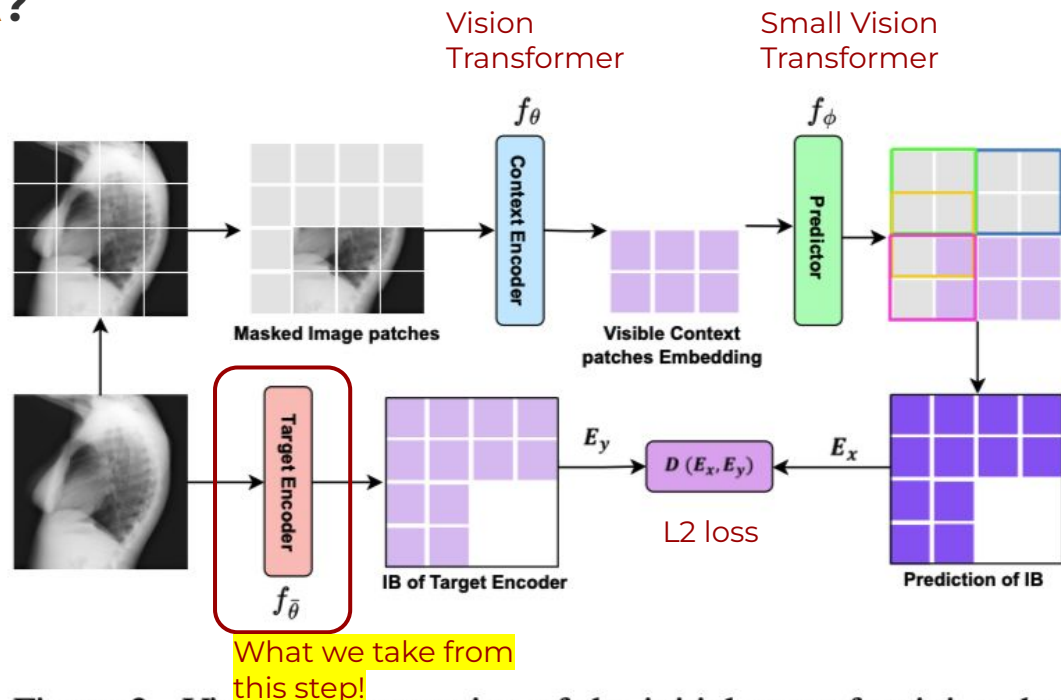





Figure 2. Visual representation of the initial step of training the encoder component of the InVERGe model.



(a) embedding space



(b) complete collapse



(c) dimensional collapse


Figure 1: Illustration of the collapsing problem. For complete collapse, the embedding vectors collapse to same point. For dimensional collapse, the embedding vectors only span a lower dimensional space.

# Cross-Modal Query Fusion Layer (CMQFL)

2/3

text-grounded image  
embeddings ( $Z$ )

32x1024



16x16=256 patches  
+ 1 cls vector = 257

257x1280



Target Encoder



projected so  
it is 257x1024

Feed Forward

Cross Attention

Bi-Self Attention

32x1024



Learnable Queries

text-grounded image embeddings ( $Z$ )

32x1024

Multi-Modality Matching

Multimodal Contrastive Learning

Image-Grounded Text Generation

16x16=256 patches  
+ 1 cls vector = 257

257x1280

projected so  
it is 257x1024

Target Encoder



Feed Forward

Cross Attention

Bi-Self Attention

Feed Forward

Cross Attention

Bi-Self Attention

32x1024

Learnable Queries

Original Report



text-grounded image embeddings ( $Z$ )

32x1024

Multi-Modality Matching

Multimodal Contrastive Learning

Image-Grounded Text Generation

Q-former  
(functions as both  
encoder and  
decoder for text)

16x16=256 patches  
+ 1 cls vector = 257

257x1280

projected so  
it is 257x1024

Target Encoder



Feed Forward

Cross Attention

Bi-Self Attention

Feed Forward

Cross Attention

Bi-Self Attention

32x1024

Learnable Queries

Original Report

# Multimodal Contrastive Learning (MCL)

Loss Term 1/3

text-grounded image embeddings ( $Z$ )

32x1024



Multimodal Contrastive Learning

16x16=256 patches  
+ 1 cls vector = 257

257x1280



projected so  
it is 257x1024

Target Encoder



Feed Forward

Cross Attention

Self Attention

32x1024



Learnable Queries

[CLS] token

Run Q-former  
as an encoder

Feed Forward

Self Attention

Original Report

# Multimodal Contrastive Learning

## What are we contrasting?

**Image Side:** After the CMQFL layer, we have:

$$Z_i = z_i^1, z_i^2, \dots, z_i^{32} \in \mathbb{R}^{(32 \times 1024)}$$

The 32 text-grounded image embeddings for image  $i$

**Text Side:** you take the [CLS] token from your BERT-based text transformer for report  $j$  call it:

$$t_j = t_{cls}^{(j)} \in \mathbb{R}^{1024}$$

A batch of size  $B$  gives you  $Z_1, \dots, Z_B$  and  $t_1, \dots, t_B$ .

# Multimodal Contrastive Learning

## What are we contrasting?

**Image Side:** After the CMQFL layer, we have:

$$Z_i = z_i^1, z_i^2, \dots, z_i^{32} \in \mathbb{R}^{(32 \times 1024)}$$

The 32 text-grounded image embeddings for image  $i$

**Text Side:** you take the [CLS] token from your BERT-based text transformer for report  $j$  call it:

$$t_j = t_{cls}^{(j)} \in \mathbb{R}^{1024}$$

A batch of size  $B$  gives you  $Z_1, \dots, Z_B$  and  $t_1, \dots, t_B$ .

## Turning them into comparable vectors:

Before you compare, you run each through small “head” projections and L2-normalize:

$$\tilde{z}_i^k = L_z(z_i^k), ||\tilde{z}_i^k||_2 = 1,$$

$$\tilde{t}_j = L_t(t_j), ||\tilde{t}_j||_2 = 1,$$

Both  $L_z$  and  $L_t$  are small linear layers mapping  $\mathbb{R}^{1024} \rightarrow \mathbb{R}^{1024}$  plus a normalization.

# Multimodal Contrastive Learning

## What are we contrasting?

**Image Side:** After the CMQFL layer, we have:

$$Z_i = z_i^1, z_i^2, \dots, z_i^{32} \in \mathbb{R}^{(32 \times 1024)}$$

The 32 text-grounded image embeddings for image  $i$

**Text Side:** you take the [CLS] token from your BERT-based text transformer for report  $j$  call it:

$$t_j = t_{cls}^{(j)} \in \mathbb{R}^{1024}$$

A batch of size  $B$  gives you  $Z_1, \dots, Z_B$  and  $t_1, \dots, t_B$ .

## Turning them into comparable vectors:

Before you compare, you run each through small “head” projections and L2-normalize:

$$\tilde{z}_i^k = L_z(z_i^k), ||\tilde{z}_i^k||_2 = 1,$$

$$\tilde{t}_j = L_t(t_j), ||\tilde{t}_j||_2 = 1,$$

Both  $L_z$  and  $L_t$  are small linear layers mapping  $\mathbb{R}^{1024} \rightarrow \mathbb{R}^{1024}$  plus a normalization.

## Defining a single image-text score:

Instead of collapsing all 32 text-grounded image embeddings by averaging, we take the single **maximum** dot-product across all embeddings:

$$sim_{text \rightarrow img}(j, i) = 1/\tau \max_{k=1 \dots 32} \langle \tilde{t}_j, \tilde{z}_i^k \rangle$$

$\tau$  is a temperature scalar that sharpens (or softens) the distribution.

We also compute the **symmetric** score for text  $\rightarrow$  image:

$$sim_{img \rightarrow text}(i, j) = 1/\tau \max_{k=1 \dots 32} \langle \tilde{z}_i^k, \tilde{t}_j \rangle$$

# Multimodal Contrastive Learning

## What are we contrasting?

**Image Side:** After the CMQFL layer, we have:

$$Z_i = z_i^1, z_i^2, \dots, z_i^{32} \in \mathbb{R}^{(32 \times 1024)}$$

The 32 text-grounded image embeddings for image  $i$

**Text Side:** you take the [CLS] token from your BERT-based text transformer for report  $j$ , call it:

$$t_j = t_{cls}^{(j)} \in \mathbb{R}^{1024}$$

A batch of size  $B$  gives you  $Z_1, \dots, Z_B$  and  $t_1, \dots, t_B$ .

## Turning them into comparable vectors:

Before you compare, you run each through small “head” projections and L2-normalize:

$$\tilde{z}_i^k = L_z(z_i^k), \|\tilde{z}_i^k\|_2 = 1,$$

$$\tilde{t}_j = L_t(t_j), \|\tilde{t}_j\|_2 = 1,$$

Both  $L_z$  and  $L_t$  are small linear layers mapping  $\mathbb{R}^{1024} \rightarrow \mathbb{R}^{1024}$  plus a normalization.



## Defining a single image-text score:

Instead of collapsing all 32 text-grounded image embeddings by averaging, we take the single **maximum** dot-product across all embeddings:

$$\text{sim}_{\text{text} \rightarrow \text{img}}(j, i) = 1/\tau \max_{k=1 \dots 32} \langle \tilde{t}_j, \tilde{z}_i^k \rangle$$

$\tau$  is a temperature scalar that sharpens (or softens) the distribution.

We also compute the **symmetric** score for text  $\rightarrow$  image:

$$\text{sim}_{\text{img} \rightarrow \text{text}}(i, j) = 1/\tau \max_{k=1 \dots 32} \langle \tilde{z}_i^k, \tilde{t}_j \rangle$$

## Building the contrastive loss:

We form two  $B \times B$  similarity matrices

$$S_{i,j}^{(i \rightarrow t)} = \text{sim}_{\text{img} \rightarrow \text{text}}(i, j) \quad S_{j,i}^{(t \rightarrow i)} = \text{sim}_{\text{text} \rightarrow \text{img}}(j, i)$$

As we know image  $i$  should match text  $i$  in the batch, we create a target vector  $Y = [0, 1, 2, \dots, B-1]$  meaning “the correct text for image  $i$  is at position  $i$  and vice versa.

Thus,

$$L_{mcl} = 1/2[CE(Y, S^{(i \rightarrow t)}) + CE(Y, S^{(t \rightarrow i)})]$$

We treat each row of  $S$  as a distribution over possible matches and push the correct one to 1.

# Masked Language Modelling (MLM)

Loss Term 2/3



text-grounded image embeddings ( $Z$ )

32x1024



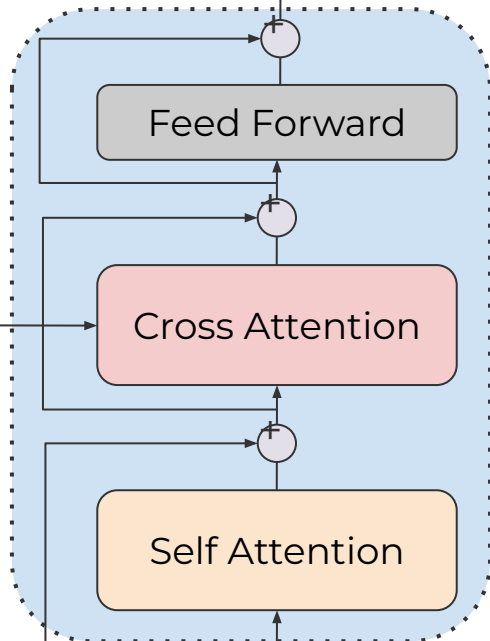
16x16=256 patches  
+ 1 cls vector = 257

257x1280



projected so  
it is 257x1024

Target Encoder



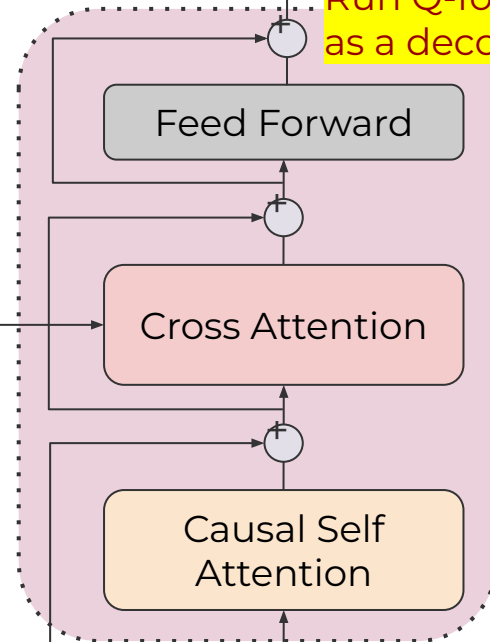
32x1024



Learnable Queries

Image-Grounded  
Text Generation

Run Q-former  
as a decoder



Original Report

## Masked Language Modelling (MLM):

1. Switch Q-former to **decoder**.
2. It uses a **causal self-attention** for text input (so text can't peek forward in the Q-former's self attention block)
3. **Cross-attention** block attends these *text representations* with KV caches from the *text-grounded image embeddings*.

The loss is then simply the **standard next-token (teacher forcing) cross-entropy** over every real token:

$$L_{MLM} = -1/N \sum_{b=1}^B \sum_{t=1}^{T-1} \mathbf{1}[y_{b,t+1} \neq -100] \log p(y_{b,t+1} | \text{context up to } t)$$


# Multi-Modality Matching (MMM)

Loss Term 3/3

text-grounded image embeddings ( $Z$ )

32x1024

Multi-Modality Matching



Binary Classifier

Logits (match, no\_match)

(0.2, 0.8), (0.4, 0.5), (0.8, 1), ...

avg

(0.5, 0.8)

Feed Forward

Cross Attention

Self Attention

Bi-Self Attention

Feed Forward

Self Attention

Original Report

16x16=256 patches  
+ 1 cls vector = 257

257x1280



Target Encoder



projected so  
it is 257x1024

32x1024



Learnable Queries

## Multi-Modality Matching (MMM):

The goal here is to teach the model to **distinguish** matched image-text pairs from hard negatives.

# Multi-Modality Matching (MMM):

The goal here is to teach the model to **distinguish** matched image-text pairs from hard negatives.

## Batch Construction

For each of the  $B$  real  $(I_b, T_b)$  pairs:

1. Positive  $(I_b, T_b)$
2. Hard-negative image  $(I_b^-, T_b)$  sampled via highest text  $\rightarrow$  image contrastive weights.
3. Hard-negative text  $(I_b, T_b^-)$  sampled via highest image  $\rightarrow$  text contrastive weights

Stack into a mega-batch of size  $3B$ .

# Multi-Modality Matching (MMM):

The goal here is to teach the model to **distinguish** matched image-text pairs from hard negatives.

## Batch Construction

For each of the  $B$  real  $(I_b, T_b)$  pairs:

1. Positive  $(I_b, T_b)$
2. Hard-negative image  $(I_b^-, T_b)$  sampled via highest text  $\rightarrow$  image contrastive weights.
3. Hard-negative text  $(I_b, T_b^-)$  sampled via highest image  $\rightarrow$  text contrastive weights

Stack into a mega-batch of size  $3B$ .

## Encoder and Bidirectional Self-Attention

We prepend the **32 learnable query tokens** so that the **bi-directional self-attention acts over [queries | text tokens]**

- *(no causal masking, instead full BERT-style attention among tokens)*

The **cross-attention mechanism** ensures the query slots attend into the image patch embeddings.

# Multi-Modality Matching (MMM):

The goal here is to teach the model to **distinguish** matched image-text pairs from hard negatives.

## Batch Construction

For each of the  $B$  real  $(I_b, T_b)$  pairs:

1. Positive  $(I_b, T_b)$
2. Hard-negative image  $(I_b^-, T_b)$  sampled via highest text  $\rightarrow$  image contrastive weights.
3. Hard-negative text  $(I_b, T_b^-)$  sampled via highest image  $\rightarrow$  text contrastive weights

Stack into a mega-batch of size  $3B$ .

## Encoder and Bidirectional Self-Attention

We prepend the **32 learnable query tokens** so that the **bi-directional self-attention acts over [queries | text tokens]**

- *(no causal masking, instead full BERT-style attention among tokens)*

The **cross-attention mechanism** ensures the query slots attend into the image patch embeddings.

## Classification Head

1. Take the final hidden states of the 32 queries  $\rightarrow$  shape  $(3B, 32, D)$
2. Apply a small **linear layer** to each query  $\rightarrow$  2-way logits ("match" vs. "non-match")  $\rightarrow (3B, 32, 2)$
3. Mean-pool across the 32 queries  $\rightarrow$  final logits  $(3B, 2)$



# Multi-Modality Matching (MMM):

The goal here is to teach the model to **distinguish** matched image-text pairs from hard negatives.

## Batch Construction

For each of the  $B$  real  $(I_b, T_b)$  pairs:

1. Positive  $(I_b, T_b)$
2. Hard-negative image  $(I_b^-, T_b)$  sampled via highest text  $\rightarrow$  image contrastive weights.
3. Hard-negative text  $(I_b, T_b^-)$  sampled via highest image  $\rightarrow$  text contrastive weights

Stack into a mega-batch of size  $3B$ .

## Encoder and Bidirectional Self-Attention

We prepend the **32 learnable query tokens** so that the **bi-directional self-attention acts over [queries | text tokens]**

- *(no causal masking, instead full BERT-style attention among tokens)*

The **cross-attention mechanism** ensures the query slots attend into the image patch embeddings.

## Classification Head

1. Take the final hidden states of the 32 queries  $\rightarrow$  shape  $(3B, 32, D)$
2. Apply a small **linear layer** to each query  $\rightarrow$  2-way logits ("match" vs. "non-match")  $\rightarrow (3B, 32, 2)$
3. Mean-pool across the 32 queries  $\rightarrow$  final logits  $(3B, 2)$

## Loss Function

**Cross-entropy** over the  $3B$  logits with labels  $y_i \in \{0, 1\}$

$$L_{MMM} = -1/(3B) \sum_{i=1}^{3B} [y_i \log p_{i,1} + (1 - y_i) \log p_{i,0}]$$

Positives get  $y_i = 1$ ;  
Both negatives get  $y_i = 0$ .

$$L_{pretraining} = L_{mcl} + L_{mlm} + L_{mmm}$$

(CMQFL Pretraining Objective)

# Decoder

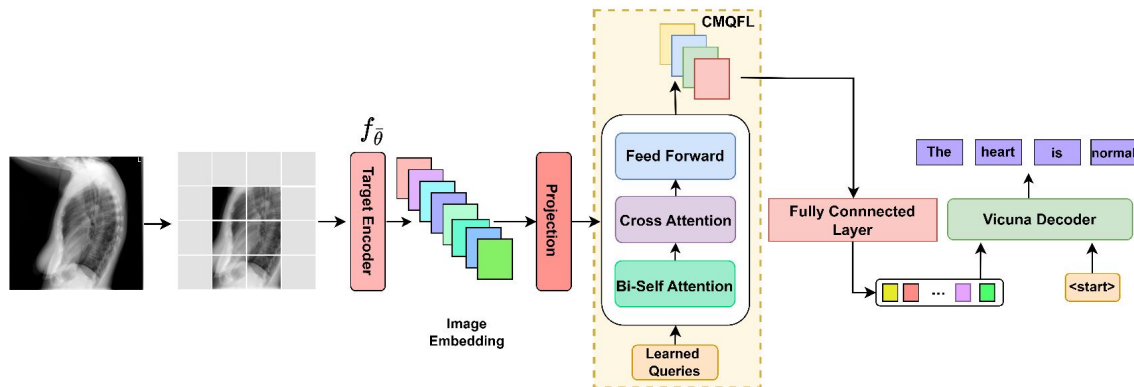
3/3

# Decoder

In the second stage (fine-tuning the VICUNA decoder):

1. **Text-grounded image embeddings** are provided as an *explicit prefix* of embedding vectors passed into VICUNA decoder
  - a. *acting as soft visual cues that condition the decoder on the visual context received by the CMQFL layer*
2. Use frozen CMQFL layer (encoder) and **exclusively employ the MLM loss.**

*“This not only eases the LLM’s burden in learning the alignment between sight and language but also reduces the problem of catastrophic forgetting.”*



# Implementation Details

# Datasets

## For Training the Image Encoder:

### 1. NIH

## For Generating the Reports:

### 1. MIMIC-CXR

- a. Images + corresponding reports
- b. We specifically focus on the 'p10' folder, which contains 36,337 images.

### 2. IU-Xray

- a. 7,470 chest X-ray images, each paired with a corresponding radiology report, totalling 3,955 reports

### 3. CDD-CESM

- a. 1003 low-energy images, featuring CC (Craniocaudal) and MLO (Mediolateral Oblique) views for both breasts. This dataset is derived from 326 female patients.

# Metrics

1. BLEU (Precision-based)
2. METEOR (Precision, Recall, Stemming, Synonyms, and Alignment)
3. ROUGE-L (Recall-based)

*Reference:* "The cat sat on the mat"  
*Our Guess:* "The cat sat on a mat"

## Unigrams:

Reference: {The, cat, sat, on, the, mat}

Candidate: {The, cat, sat, on, a, mat}

**BLEU-1 = % ~ 0.83**

*Reference:* "The cat sat on the mat"  
*Our Guess:* "A feline was sitting on a rug"

## Aligns:

"cat" <-> "feline" (synonym)

"sat" <-> "sitting" (stem)

"mat" <-> "rug" (synonym)

**METEOR ~ 0.5 or higher**

*Reference:* "The cat sat on the mat"

*Our Guess:* "The cat is sitting on the mat"

*LCS:* "The cat on the mat"

|LCS| = 5 words

Reference length = 6 words

**ROUGE-L = % ~ 0.83**

BLEU and METEOR were initially developed for *assessing machine translation quality*, while ROUGE-L is specifically designed for *evaluating the quality of textual summaries*.

# Results



Table 1. Comparison of the proposed InVERGe and other SOTA methods on the MIMIC-CXR, IU and CDD-CESM datasets. A higher value indicates superior performance in all categories. The best performance is highlighted in **bold**.

DataSet	Models	BLEU-1	BLEU-2	BLEU-3	BLEU-4	METEOR	ROUGE-L
MIMIC-CXR	Show-Tell [48]	0.279	0.159	0.101	0.076	0.121	0.251
	AdaAtt [35]	0.28	0.16	0.105	0.079	0.122	0.254
	Attn2in [40]	0.331	0.213	0.121	0.081	0.131	0.271
	Up-Down [3]	0.318	0.191	0.113	0.071	0.128	0.267
	R2Gen [13]	0.311	0.186	0.112	0.077	0.125	0.265
	M2 Transformer [16]	0.347	0.211	0.122	0.085	0.140	0.269
	X-REM [23]	0.314	0.188	0.112	0.069	0.121	0.266
	BLIP-2 [32]	0.377	0.221	0.125	0.088	0.152	0.274
	R2GenGPT (Deep) [52]	0.392	0.229	0.129	<b>0.101</b>	0.159	0.283
	Proposed (InVERGe)	<b>0.425</b>	<b>0.240</b>	<b>0.132</b>	0.100	<b>0.175</b>	<b>0.309</b>
IU X-Ray	Show-Tell [48]	0.341	0.203	0.140	0.079	0.123	0.321
	AdaAtt [35]	0.434	0.283	0.195	0.126	0.143	0.342
	Attn2in [40]	0.402	0.262	0.188	0.119	0.133	0.338
	Up-Down [3]	0.383	0.248	0.176	0.116	0.129	0.337
	R2Gen [13]	0.421	0.261	0.170	0.121	0.139	0.335
	M2 Transformer [16]	0.456	0.312	0.202	0.151	0.168	0.351
	X-REM [23]	0.426	0.263	0.171	0.119	0.135	0.341
	BLIP-2 [32]	0.476	0.273	0.210	0.168	0.181	0.372
	R2GenGPT (Deep) [52]	0.481	0.301	0.214	<b>0.169</b>	0.189	0.375
	Proposed (InVERGe)	<b>0.499</b>	<b>0.324</b>	<b>0.226</b>	0.168	<b>0.195</b>	<b>0.384</b>
CDD-CESM	Show-Tell [48]	0.284	0.165	0.109	0.079	0.153	0.264
	AdaAtt [35]	0.293	0.169	0.116	0.080	0.180	0.269
	Attn2in [40]	0.340	0.217	0.124	0.081	0.240	0.306
	Up-Down [3]	0.338	0.204	0.123	0.079	0.237	0.310
	R2Gen [13]	0.335	0.199	0.122	0.077	0.213	0.299
	M2 Transformer [16]	0.357	0.221	0.125	0.085	0.256	0.315
	X-REM [23]	0.333	0.197	0.119	0.074	0.210	0.297
	BLIP-2 [32]	0.382	0.235	0.139	0.102	0.301	0.342
	R2GenGPT (Deep) [52]	0.417	0.249	0.165	0.129	0.354	0.377
	Proposed (InVERGe)	<b>0.453</b>	<b>0.267</b>	<b>0.185</b>	<b>0.134</b>	<b>0.391</b>	<b>0.430</b>

Table 2. Baseline refers to one normal ViT encoder and decoder only. CPE (Context Pixel encoder) stands for our trained Model's Target Encoder. For this, we use the MIMIC-CXR dataset. Here we only use MLM objective function to check the performance.

Model	BLEU-2	METEOR	Rouge-L
Baseline	0.161	0.124	0.255
MAE + Decoder	0.178	0.060	0.208
CPE + Decoder	0.183	0.117	0.260
CPE + CMQFL + Decoder	<b>0.227</b>	<b>0.163</b>	<b>0.290</b>

Table 3. Evaluation after adding additional objective to train the CMQFL layer. For this, we use MIMIC-CXR dataset.

Model	BLEU-1	BLEU-2	METEOR	Rouge-L
MLM	0.410	0.227	0.163	0.290
MLM+MMM	0.416	0.231	0.166	0.30
MLM+MMM+MCL (InVERGe)	<b>0.425</b>	<b>0.24</b>	<b>0.175</b>	<b>0.309</b>

The combined effect of these three objective functions leads to a substantial improvement in the performance of the pre-trained model, enhancing its ability to understand and leverage both text and image data effectively.

By adding these objectives, the model's overall capabilities are noticeably enhanced. The addition of MCL focuses on strengthening the alignment between images and text, maximizing their common information. This contrastive approach encourages the model to distinguish between positive image-text pairs and negative pairs, resulting in a stronger understanding of the relationship between visual and text data. At the same time, the introduction of MMM further refines the model's cross-modal capabilities, helping to fine-tune the CMQFL layer by emphasizing alignment between images and text through matching objectives.

## JEPA-based image encoder

### Problem

**[DONE]** Models have been observed to produce **unnatural language outputs** when trained solely on raw image-text pairs

**[DONE]** High computational demand of end-to-end model training

The introduction of an advanced encoder for extracting detailed visual features enhances the model's ability to generate reports without requiring additional annotations or external task-specific knowledge. The CMQFL layer, with its three objectives, contributes to creating small yet highly informative image embeddings, promoting a more grounded vision and language representation. This approach not only improves the accuracy of the model within a shorter training period but also surpasses previous SOTA models on publicly available datasets, delivering detailed radiology reports and marking a significant advancement in the field.

Bridging  
modalities

## Paper Reference

A. Deria, K. Kumar, S. Chakraborty, D. Mahapatra and S. Roy, "InVERGe: Intelligent Visual Encoder for Bridging Modalities in Report Generation," 2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW), Seattle, WA, USA, 2024, pp. 2028-2038, doi: 10.1109/CVPRW63382.2024.00208. keywords: {Training;Visualization;Computational modeling;Radiology;Transformers;Feature extraction;Decoding;Computer Vision;Deep Learning;Medical Imaging;Report Generation;Large Language Model},